

Spring 2023: CS5710 – Machine Learning

In-Class Programming Assignment-5

Name: Chetan Kumar Reddy Naga

ID: 700743408

Link for Github:

<https://github.com/ChetanNaga/Machine-Learning.git>

Imported all the required libraries such as pandas, numpy, sklearn and etc.
And loaded the dataset “CC.csv”

```
In [1]: # importing required libraries for assignment 5 here
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
from sklearn import preprocessing, metrics
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
sns.set(style="white", color_codes=True)
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: # Principal Component Analysis
# a. Apply PCA on CC dataset.
# b. Apply k-means algorithm on the PCA result and report your observation if the silhouette score
# has improved or not?
# c. Perform Scaling+PCA+K-Means and report performance.
```

```
In [5]: dataset_CC = pd.read_csv('datasets//CC GENERAL.csv')
dataset_CC.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8950 entries, 0 to 8949
Data columns (total 18 columns):
#   Column              Non-Null Count  Dtype  
---  --
0   CUST_ID              8950 non-null   object  
1   RAI ANCF             8950 non-null   float64
```

To apply PCA we need to remove all the null values that are present in the dataset. So, removed all the null values and applied PCA on the dataset.

The screenshot shows a Jupyter Notebook window titled "Assignment5" with a last checkpoint 8 hours ago. The interface includes a menu bar (File, Edit, View, Insert, Cell, Kernel, Widgets, Help) and a toolbar with icons for file operations, running, and code execution. The notebook contains two code cells. The first cell, labeled "In [5]:", executes the command `dataset_CC.isnull().any()`, resulting in an output labeled "Out[5]:". The output is a series of attribute names followed by their null status: `CUST_ID` (False), `BALANCE` (False), `BALANCE_FREQUENCY` (False), `PURCHASES` (False), `ONEOFF_PURCHASES` (False), `INSTALLMENTS_PURCHASES` (False), `CASH_ADVANCE` (False), `PURCHASES_FREQUENCY` (False), `ONEOFF_PURCHASES_FREQUENCY` (False), `PURCHASES_INSTALLMENTS_FREQUENCY` (False), `CASH_ADVANCE_FREQUENCY` (False), `CASH_ADVANCE_TRX` (False), `PURCHASES_TRX` (False), `CREDIT_LIMIT` (True), `PAYMENTS` (False), `MINIMUM_PAYMENTS` (True), `PRC_FULL_PAYMENT` (False), `TENURE` (False), and `dtype: bool`. The second cell, labeled "In [6]:", executes the command `dataset_CC.fillna(dataset_CC.mean(), inplace=True)` followed by `dataset_CC.isnull().any()`. The output for this cell is partially visible, showing `CUST_ID` (False) and `BALANCE` (False).

```
In [5]: dataset_CC.isnull().any()

Out[5]: CUST_ID                False
        BALANCE                False
        BALANCE_FREQUENCY      False
        PURCHASES              False
        ONEOFF_PURCHASES       False
        INSTALLMENTS_PURCHASES False
        CASH_ADVANCE           False
        PURCHASES_FREQUENCY    False
        ONEOFF_PURCHASES_FREQUENCY False
        PURCHASES_INSTALLMENTS_FREQUENCY False
        CASH_ADVANCE_FREQUENCY False
        CASH_ADVANCE_TRX       False
        PURCHASES_TRX          False
        CREDIT_LIMIT           True
        PAYMENTS               False
        MINIMUM_PAYMENTS       True
        PRC_FULL_PAYMENT       False
        TENURE                 False
        dtype: bool

In [6]: dataset_CC.fillna(dataset_CC.mean(), inplace=True)
        dataset_CC.isnull().any()

Out[6]: CUST_ID                False
        BALANCE                False
```

First, Applied the elbow method to find the clusters required for performing the k-means implementation. Which is 3 from the obtained output.

On applying the K-means algorithm for the PCA result we get the Silhoutte score as 0.5109

```
score = metrics.silhouette_score(X, y_cluster_kmeans)
print("Silhouette Score: ",score)

"""
Silhouette Score- ranges from -1 to +1 , a high value indicates that the object is well matched to its own cluster and poorly
"""
```

	precision	recall	f1-score	support
0	0.00	1.00	0.00	0.0
1	0.00	1.00	0.00	0.0
2	0.00	1.00	0.00	0.0
6	1.00	0.00	0.00	204.0
7	1.00	0.00	0.00	190.0
8	1.00	0.00	0.00	196.0
9	1.00	0.00	0.00	175.0
10	1.00	0.00	0.00	236.0
11	1.00	0.00	0.00	365.0
12	1.00	0.00	0.00	7584.0
accuracy			0.00	8950.0
macro avg	0.70	0.30	0.00	8950.0
weighted avg	1.00	0.00	0.00	8950.0

```
[[ 0 0 0 0 0 0 0 0 0 0 0]
 [ 0 0 0 0 0 0 0 0 0 0 0]
 [ 0 0 0 0 0 0 0 0 0 0 0]
 [175 1 28 0 0 0 0 0 0 0 0]
 [173 2 15 0 0 0 0 0 0 0 0]
 [169 0 27 0 0 0 0 0 0 0 0]
 [149 0 26 0 0 0 0 0 0 0 0]
 [188 1 47 0 0 0 0 0 0 0 0]
 [284 3 78 0 0 0 0 0 0 0 0]
 [5389 126 2069 0 0 0 0 0 0 0 0]]

Accuracy for our Training dataset with PCA: 0.0
Silhouette Score: 0.5109307274319468
```

Reload the dataset to perform the scaling

```
In [11]: nclusters = 3 # this is the k in kmeans
km = KMeans(n_clusters=nclusters)
km.fit(X)

# predict the cluster for each data point
y_cluster_kmeans = km.predict(X)

# Summary of the predictions made by the classifier
print(classification_report(y, y_cluster_kmeans, zero_division=1))
print(confusion_matrix(y, y_cluster_kmeans))

train_accuracy = accuracy_score(y, y_cluster_kmeans)
print("\nAccuracy for our Training dataset with PCA:", train_accuracy)

#Calculate silhouette Score
score = metrics.silhouette_score(X, y_cluster_kmeans)
print("Silhouette Score: ",score)

"""
Silhouette Score- ranges from -1 to +1 , a high value indicates that the object is well matched to its own cluster and poorly
"""
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Apply the Scaling and perform PCA on the dataset

```
In [13]: #Scaling
scaler = StandardScaler()
scaler.fit(x)
X_scaled_array = scaler.transform(x)
#PCA
pca = PCA(3)
x_pca = pca.fit_transform(X_scaled_array)
principalDf = pd.DataFrame(data = x_pca, columns = ['principal component 1', 'principal component 2', 'principal component 3'])
finalDf = pd.concat([principalDf, dataset_CC.iloc[:, -1]], axis = 1)
finalDf.head()
```

```
Out[13]:
```

	principal component 1	principal component 2	principal component 3	TENURE
0	-1.718893	-1.072939	0.535670	12
1	-1.169306	2.509320	0.628027	12
2	0.938414	-0.382600	0.161198	12
3	-0.907503	0.045859	1.521689	12
4	-1.637830	-0.684975	0.425658	12

```
In [14]: X = finalDf.iloc[:, 0:-1]
y = finalDf["TENURE"]
print(X.shape, y.shape)

(8950, 3) (8950,)
```

```
In [15]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.34, random_state=0)
```

Now performing the k-means on the scaled PCA data, which gives the result of 0.383. Which has reduced from the previous k-means value.

```
"""
Silhouette Score- ranges from -1 to +1 , a high value indicates that the object is well matched to its own cluster and poorly
"""
```

	precision	recall	f1-score	support
0	0.00	1.00	0.00	0.0
1	0.00	1.00	0.00	0.0
2	0.00	1.00	0.00	0.0
6	1.00	0.00	0.00	65.0
7	1.00	0.00	0.00	55.0
8	1.00	0.00	0.00	68.0
9	1.00	0.00	0.00	57.0
10	1.00	0.00	0.00	85.0
11	1.00	0.00	0.00	103.0
12	1.00	0.00	0.00	2610.0
accuracy			0.00	3043.0
macro avg	0.70	0.30	0.00	3043.0
weighted avg	1.00	0.00	0.00	3043.0

```
[[ 0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0]
 [ 41 21  3  0  0  0  0  0  0  0]
 [ 42 12  1  0  0  0  0  0  0  0]
 [ 57 10  1  0  0  0  0  0  0  0]
 [ 35 22  0  0  0  0  0  0  0  0]
 [ 63 17  5  0  0  0  0  0  0  0]
 [ 69 30  4  0  0  0  0  0  0  0]
 [1763 450 397  0  0  0  0  0  0  0]]
```

Accuracy for our Training dataset with PCA: 0.0
Silhouette Score: 0.383322340968964

Load the pd_speech_features dataset

```
In [18]: dataset_pd = pd.read_csv('datasets//pd_speech_features.csv')
dataset_pd.info()

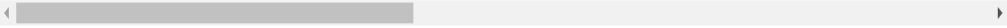
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 756 entries, 0 to 755
Columns: 755 entries, id to class
dtypes: float64(749), int64(6)
memory usage: 4.4 MB

In [19]: dataset_pd.head()

Out[19]:
```

	id	gender	PPE	DFA	RPDE	numPulses	numPeriodsPulses	meanPeriodPulses	stdDevPeriodPulses	locPctJitter	...	tqwt_kurtosisValue_dec_20
0	0	1	0.85247	0.71826	0.57227	240	239	0.008064	0.000087	0.00218	...	1.562
1	0	1	0.76686	0.69481	0.53966	234	233	0.008258	0.000073	0.00195	...	1.558
2	0	1	0.85083	0.67604	0.58982	232	231	0.008340	0.000060	0.00176	...	1.564
3	1	0	0.41121	0.79672	0.59257	178	177	0.010858	0.000183	0.00419	...	3.780
4	1	0	0.32790	0.79782	0.53028	236	235	0.008162	0.002669	0.00535	...	6.172

5 rows × 755 columns



```
In [20]: dataset_pd.isnull().any()
```

We are doing the scaling on the pd_speech_features dataset and then apply PCA with k=3 value

```
X_Scale = scaler.fit_transform(X)
```

In [23]:

```
# Apply PCA with k =3
pca3 = PCA(n_components=3)
principalComponents = pca3.fit_transform(X_Scale)

principalDf = pd.DataFrame(data = principalComponents, columns = ['principal component 1', 'principal component 2', 'Principal Component 3', 'class'])

finalDf = pd.concat([principalDf, dataset_pd[['class']]], axis = 1)
finalDf.head()
```

Out[23]:

	principal component 1	principal component 2	Principal Component 3	class
0	-10.047372	1.471076	-6.846402	1
1	-10.637725	1.583749	-6.830976	1
2	-13.516185	-1.253542	-6.818696	1
3	-9.155084	8.833601	15.290906	1
4	-6.764470	4.611468	15.637121	1

In [24]:

```
X = finalDf.drop('class',axis=1).values
y = finalDf['class'].values
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.34,random_state=0)
```

We have used the SVM classifier to report the performance which is 0.810 and got the classification report

```
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.34,random_state=0)
```

In [25]:

```
#2.c Support Vector Machine's
from sklearn.svm import SVC

svmClassifier = SVC()
svmClassifier.fit(X_train, y_train)

y_pred = svmClassifier.predict(X_test)

# Summary of the predictions made by the classifier
print(classification_report(y_test, y_pred, zero_division=1))
print(confusion_matrix(y_test, y_pred))
# Accuracy score
glass_acc_svc = accuracy_score(y_pred,y_test)
print('accuracy is',glass_acc_svc )

#Calculate sihouette Score
score = metrics.silhouette_score(X_test, y_pred)
print("Sihouette Score: ",score)
```

	precision	recall	f1-score	support
0	0.67	0.42	0.51	62
1	0.84	0.93	0.88	196
accuracy			0.81	258
macro avg	0.75	0.68	0.70	258
weighted avg	0.80	0.81	0.79	258

```
[[ 26 36]
 [ 13 183]]
accuracy is 0.810077519379845
Sihouette Score: 0.2504463929631217
```

Load the Iris dataset to perform LDA.

Apply Standard Scaling on the data

```
sinusoidal score: 0.230440325031217

In [26]: #3.Apply Linear Discriminant Analysis (LDA) on Iris.csv dataset to reduce dimensionality of data to k=2.
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
dataset_iris = pd.read_csv('datasets//Iris.csv')
dataset_iris.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype  
---  -
0    Id              150 non-null   int64   
1    SepalLengthCm   150 non-null   float64  
2    SepalWidthCm    150 non-null   float64  
3    PetalLengthCm   150 non-null   float64  
4    PetalWidthCm    150 non-null   float64  
5    Species         150 non-null   object  
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

```
In [27]: dataset_iris.isnull().any()
```

```
Out[27]: Id              False
SepalLengthCm          False
SepalWidthCm           False
PetalLengthCm          False
PetalWidthCm           False
Species                False
```

Reducing the dataset dimensionality to k=2

```
jupyter Assignment5 Last Checkpoint: 9 hours ago (autosaved) Logout
File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel)
In [27]: dataset_iris.isnull().any()
Out[27]: Id False
SepallengthCm False
SepalWidthCm False
PetalLengthCm False
PetalWidthCm False
Species False
dtype: bool

In [28]: x = dataset_iris.iloc[:,1:-1]
y = dataset_iris.iloc[:, -1]
print(x.shape,y.shape)
(150, 4) (150,)

In [29]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=0)

In [30]: sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
le = LabelEncoder()
y = le.fit_transform(y)

In [31]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
lda = LDA(n_components=2)
X_train = lda.fit_transform(X_train, y_train)
X_test = lda.transform(X_test)
print(X_train.shape,X_test.shape)
(105, 2) (45, 2)
```

4. Briefly identify the difference between PCA and LDA?

PCA performs better in case where number of samples per class is less. Whereas LDA works better with large dataset having multiple classes; class separability is an important factor while reducing dimensionality. PCA finds directions of maximum variance regardless of class labels while LDA finds directions of maximum class separability.

